

# Elasticity Detection: A Building Block for Internet Congestion Control[4]


King-Siong Si

Institute of Multimedia Knowledge Fusion and Engineering,  
Xi'an JiaoTong University  
*sjsinx@stu.xjtu.edu.cn*

January 24, 2024



- ① Background
- ② Methodology
- ③ Experiment
- ④ Other Details
- ⑤ References




# ① Background

## ② Methodology

## ③ Experiment

## ④ Other Details

## ⑤ References



# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput



# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.




# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.




# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.



# To Be, Or Not to Be?


- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay






# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.




# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.




# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.
  - **High queueing delay** even when cross-traffic is not aggressive



# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.
  - **High queueing delay** even when cross-traffic is not aggressive
- Information about cross traffic can improve CC.




# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.
  - **High queueing delay** even when cross-traffic is not aggressive
- Information about cross traffic can improve CC.
  - Adjust CC behavior based on the nature of cross traffic.

# To Be, Or Not to Be?


- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.
  - **High queueing delay** even when cross-traffic is not aggressive
- Information about cross traffic can improve CC.
  - Adjust CC behavior based on the nature of cross traffic.
  - Use Cubic when cross traffic is competing.



# To Be, Or Not to Be?

- Delay-Based Algorithms: Low Throughput
  - Examples: Vegas[2], Copa[1], etc.
  - Achieve lower delays.
  - **Poor throughput** against deployed algorithms like Cubic[5], NewReno[6], BBR[3], etc.
- Deployed Algorithms: High Delay
  - Examples: Cubic, NewReno, etc.
  - Compete well for throughput.
  - **High queueing delay** even when cross-traffic is not aggressive
- Information about cross traffic can improve CC.
  - Adjust CC behavior based on the nature of cross traffic.
  - Use Cubic when cross traffic is competing.
  - Use delay-based CC algorithms otherwise.

- 1 Background
- 2 Methodology**
- 3 Experiment
- 4 Other Details
- 5 References





# Questions

- Question 1: What kind of information?



# Questions

- Question 1: What kind of information?
- Elasticity




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.



# Questions


- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.
- Otherwise,






# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.
- Otherwise,
  - Use TCP-competitive mode, like Cubic, BBR, etc.




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.
- Otherwise,
  - Use TCP-competitive mode, like Cubic, BBR, etc.
- Switch between two modes to achieve




# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.
- Otherwise,
  - Use TCP-competitive mode, like Cubic, BBR, etc.
- Switch between two modes to achieve
  - Low delays



# Questions

- Question 1: What kind of information?
- Elasticity
- Elastic cross traffic: If contains *any* elastic flow.
  - Elastic flow: Adjust sending rate in response to available bandwidth.
  - Inelastic flow: Otherwise.
- If cross traffic is inelastic then,
  - Use delay-controlling mode, like Vegas, Copa, etc.
- Otherwise,
  - Use TCP-competitive mode, like Cubic, BBR, etc.
- Switch between two modes to achieve
  - Low delays
  - Same throughput as the deployed TCP-competitive CC algorithms



# Questions

- Question 2: How to detect elasticity?



# Questions

- Question 2: How to detect elasticity?
- Strawman: Compare total queueing delay and “self-inflicted” delay.




Figure 1: Instantaneous delay measurements do not reveal elasticity.<sup>1</sup>

<sup>1</sup>From Figure 2 in [4].

# Questions

- Question 2: How to detect elasticity?
- Strawman: Compare total queueing delay and “self-inflicted” delay.




Figure 1: Instantaneous delay measurements do not reveal elasticity.<sup>1</sup>

- Nimbus

<sup>1</sup>From Figure 2 in [4].

# Questions

- Question 2: How to detect elasticity?
- Strawman: Compare total queueing delay and “self-inflicted” delay.




Figure 1: Instantaneous delay measurements do not reveal elasticity.<sup>1</sup>

- Nimbus
  - Modulate sending rate to create variations in available bandwidth.

<sup>1</sup>From Figure 2 in [4].



# Questions

- Question 2: How to detect elasticity?
- Strawman: Compare total queueing delay and “self-inflicted” delay.




Figure 1: Instantaneous delay measurements do not reveal elasticity.<sup>1</sup>

- Nimbus
  - Modulate sending rate to create variations in available bandwidth.
  - Monitor how the cross traffic responds to these variations.

<sup>1</sup>From Figure 2 in [4].


# Steps

- Step 1: Measure rate of the cross traffic.




# Steps

- Step 1: Measure rate of the cross traffic.




# Steps

- Step 1: Measure rate of the cross traffic.



# Steps


- Step 1: Measure rate of the cross traffic.



- Input:  $\frac{S(t)}{S(t) + z(t)}$ ,

# Steps


- Step 1: Measure rate of the cross traffic.



- Input:  $\frac{S(t)}{S(t)+z(t)}$ , Output:  $\frac{R(t)}{\mu}$ ,

# Steps


- Step 1: Measure rate of the cross traffic.



- Input:  $\frac{S(t)}{S(t)+z(t)}$ , Output:  $\frac{R(t)}{\mu}$ ,  $\Rightarrow \frac{S(t)}{S(t)+z(t)} = \frac{R(t)}{\mu}$

# Steps

- Step 1: Measure rate of the cross traffic.



- Input:  $\frac{S(t)}{S(t)+z(t)}$ , Output:  $\frac{R(t)}{\mu}$ ,  $\Rightarrow \frac{S(t)}{S(t)+z(t)} = \frac{R(t)}{\mu}$
- $\Rightarrow \hat{z}(t) = \mu \frac{S(t)}{R(t)} - S(t)$



# Steps

- Step 2: Create variations in available bandwidth.



# Steps

- Step 2: Create variations in available bandwidth.
- Send packets in pulses.



# Steps

- Step 2: Create variations in available bandwidth.
- Send packets in pulses.
- $S(t) = S_{expected}(t) + A \sin(2\pi ft)$ .



# Steps

- Step 2: Create variations in available bandwidth.
- Send packets in pulses.
- $S(t) = S_{expected}(t) + A \sin(2\pi ft)$ .




Figure 2: Cross traffic's reaction to pulses.<sup>1</sup>

<sup>1</sup>From Figure 3 in [4].


# Steps

- Step 3: Classify cross traffic elasticity.




# Steps

- Step 3: Classify cross traffic elasticity.
- Measure the *cross-correlation* between  $S(t)$  and  $z(t)$ .




# Steps

- Step 3: Classify cross traffic elasticity.
- Measure the *cross-correlation* between  $S(t)$  and  $z(t)$ . (x)
- Align  $S(t)$  and  $z(t)$  is challenging.



# Steps

- Step 3: Classify cross traffic elasticity.
- Measure the *cross-correlation* between  $S(t)$  and  $z(t)$ . (x)
- Align  $S(t)$  and  $z(t)$  is challenging.





# Steps

- Step 3: Classify cross traffic elasticity.
- Measure the *cross-correlation* between  $S(t)$  and  $z(t)$ . (x)
- Align  $S(t)$  and  $z(t)$  is challenging.



# Steps

- Step 3: Classify cross traffic elasticity.



# Steps

- Step 3: Classify cross traffic elasticity.
- From time to frequency domain



# Steps

- Step 3: Classify cross traffic elasticity.
- From time to frequency domain




Figure 3: Cross traffic FFT for elastic and inelastic traffic.<sup>1</sup>

<sup>1</sup>From Figure 4 in [4].

# Steps

- Step 3: Classify cross traffic elasticity.
- From time to frequency domain




Figure 3: Cross traffic FFT for elastic and inelastic traffic.<sup>1</sup>

- Elastic cross traffic exhibits a pronounced peak at  $f_p$  compared to the neighboring frequencies.

<sup>1</sup>From Figure 4 in [4].

# Steps

- Step 3: Classify cross traffic elasticity.
- From time to frequency domain




Figure 3: Cross traffic FFT for elastic and inelastic traffic.<sup>1</sup>

- Elastic cross traffic exhibits a pronounced peak at  $f_p$  compared to the neighboring frequencies.
- No need to know RTT of cross traffic.

<sup>1</sup>From Figure 4 in [4].


① Background

② Methodology

③ Experiment

④ Other Details

⑤ References



# Experiment

- Setup
  - Bottleneck rate: 96 Mbit/s,
  - Minimum RTT: 50 ms
  - Buffer: 2 BDP





# Experiment




Figure 4: Performance of NimbusCC.<sup>1</sup>

<sup>1</sup>From Figure 7 in [4].


① Background

② Methodology

③ Experiment


④ Other Details

⑤ References




## Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).




## Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.




## Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,




# Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and




# Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and
  - $f_{pd}$  in delay-control mode.



## Multiple NimbusCC Flows


- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and
  - $f_{pd}$  in delay-control mode.
- A watcher: Infers whether the pulser is pulsing at frequency  $f_{pc}$  or frequency  $f_{pd}$  by





# Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and
  - $f_{pd}$  in delay-control mode.
- A watcher: Infers whether the pulser is pulsing at frequency  $f_{pc}$  or frequency  $f_{pd}$  by
  - computing the FFT of **its receive rate**,  $R$ , at these two frequencies.




# Multiple NimbusCC Flows


- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and
  - $f_{pd}$  in delay-control mode.
- A watcher: Infers whether the pulser is pulsing at frequency  $f_{pc}$  or frequency  $f_{pd}$  by
  - computing the FFT of **its receive rate**,  $R$ , at these two frequencies.
- Pulser election: A distributed and randomized election decides which flow is the pulser and which are watchers.

# Multiple NimbusCC Flows

- NimbusCC fails if all the NimbusCC flows pulse at the same frequency( $f_p$ ).
- Set the pulser and the watchers.
- The pulser: Use two different frequencies,
  - $f_{pc}$  in TCP-competitive mode, and
  - $f_{pd}$  in delay-control mode.
- A watcher: Infers whether the pulser is pulsing at frequency  $f_{pc}$  or frequency  $f_{pd}$  by
  - computing the FFT of **its receive rate**,  $R$ , at these two frequencies.
- Pulser election: A distributed and randomized election decides which flow is the pulser and which are watchers.
  - Similar to CSMA



- ① Background
- ② Methodology
- ③ Experiment
- ④ Other Details
- ⑤ References



[1] Venkat Arun and Hari Balakrishnan.

Copa: Practical {Delay-Based} congestion control for the internet.

*In 15th USENIX Symposium on Networked Systems Design and Implementation (NSDI 18), pages 329–342, 2018.*

[2] Lawrence S Brakmo, Sean W O'malley, and Larry L Peterson.

Tcp vegas: New techniques for congestion detection and avoidance.

*In Proceedings of the conference on Communications architectures, protocols and applications, pages 24–35, 1994.*

[3] Neal Cardwell, Yuchung Cheng, C Stephen Gunn, Soheil Hassas Yeganeh, and Van Jacobson.

Bbr: Congestion-based congestion control.

*Communications of the ACM, 60(2):58–66, 2017.*

- [4] Prateesh Goyal, Akshay Narayan, Frank Cangialosi, Srinivas Narayana, Mohammad Alizadeh, and Hari Balakrishnan.  
Elasticity detection: A building block for internet congestion control.


*In Proceedings of the ACM SIGCOMM 2022 Conference*, pages 158–176, 2022.

- [5] Sangtae Ha, Injong Rhee, and Lisong Xu.  
Cubic: a new tcp-friendly high-speed tcp variant.

*ACM SIGOPS operating systems review*, 42(5):64–74, 2008.

- [6] Janey C Hoe.  
Improving the start-up behavior of a congestion control scheme for tcp.

*ACM SIGCOMM Computer Communication Review*, 26(4):270–280, 1996.



# THANKS!